Exploring window-based approaches for genotype-environment association studies

Tom R. Booker[*], **Samuel Yeaman**[†] **and Michael C. Whitlock**[*]

[*]University of British Columbia, [†]University of Calgary

1 **ABSTRACT** *I'm just using the GENETICS template because it looks nice!*

2 Here is a really concise and nicely written summary of the paper, highlighting the main findings and take home messages.

3

4 **KEYWORDS** Local Adaptation, Population Genetics, Environmental Genomics

Introduction

Alleles subject to spatially varying selection pressures may vary in frequency across a species' range in response to changing environmental conditions. For that reason, genetic variants that exhibit strong correlations with putatively selective features of the environment are often interpreted as a signature of local adaptation. Genotype-environment association (GEA) studies examine such correlations. Typically, allele frequencies are estimated in multiple locations across a species range for many genetic markers, typically single nucleotide polymorphisms (hereafter SNPs). Correlations between allele frequency and environmental variables are calculated then contrasted for sites across the genome. In the context of GEA studies, the term environment may refer to any abiotic or biotic variable that the species of interest could conceivably be adapting/adapted to.

Within a species' range, there may be wide environmental variation and such variation could potentially be correlated with patterns of gene flow or historical demography. For example, consider a hypothetical species inhabiting a large latitudinal range. If the hypothetical species had restricted migration and exhibited isolation-by-distance, neutral alleles may happen to be correlated with environmental variables that correlate with latitude simply due to population structure. For that reason, any attempt to identify loci involved in adaptation may then be stymied by an underlying correlation between presumed selection pressures and directions of gene flow; Sohail et al (2018) and Berg et al (2018) provide a clear example of this problem in an analysis of selection on human height. GEA approaches may correct for population demography when calculating correlations with the environment. For example, the commonly used BayEnv (Coop et al 2013) and BayPass (Gautier 2016) packages estimate a population covariance matrix (Ω) from SNP data, then use it as a fixed parameter when estimating correlations between the frequencies of individual SNPs and environmental variation. However, estimation of Ω from alleles with patterns of allele frequency that have been shaped by selection may obscure signals of selection. Indeed, Lotterhos (2019) analysed simulated data and found that GEA methods that corrected for population structure using Ω had less power than a comparatively simple Spearman's rank correlation between allele frequency and environment.

Closely linked SNPs are not independently inherited and have correlated evolutionary histories. Theoretical studies of local adaptation suggest that we should expect regions of the genome subject to spatially varying selection pressures to exhibit elevated linkage disequilibrium relative to the genomic background and there are at least two possible reasons why this might be the case. Firstly, under local adaptation, alleles are subject to positive selection in some parts of a species' range, but not in others. Non-beneficial genetic variants

introduced to local populations via gene flow may be removed, with a result being a build up of LD between selected alleles and linked neutral sites. This process can be thought of as a local barrier to gene flow (Barton and Bengston et al). Secondly, there is a selective advantage for alleles that are involved in local adaptation to aggregate in regions of low recombination so favourable combinations of alleles may be bound together into regions of high LD (Rieseberg 2001; Noor et al 2001; Kirkpatrick and Barton 2006; Yeaman 2012). For example, in sunflowers and *Littorina* marine snails there is evidence that regions of suppressed recombination cause alleles involved in local adaptation to be inherited together (Morales et al 2019; Todesco et al 2020) and in conifers many of the genes with the strongest signals of local adaptation are in LD with each other (Yeaman et al 2016). Of course, the two processes we have outlined are not mutually exclusive. Overall, genomic regions containing strongly selected alleles that contribute to local adaptation may potentially exhibit GEA signals at multiple linked sites.

If local adaptation induces LD and causes multiple linked sites to exhibit similar Indeed, to analyse data from Interior spruce and lodgepole pine, Yeaman et al (2016) developed a method they called the "top-candidate" test for combining information across sites in GEA studies. In this study, we perform simulation of local adaptation in genomes with linkage in order to test the performance of methods to Linked sites do not evolve independently. If the rate of recombination is low relative to the rate of migration, there may be strong autocorrelation in the coalescent histories among tightly linked sites. Under that assumption, all of the neutral SNPs present within an appropriately sized region provide independent tests of the following hypothesis, "is the genetic variation in this genomic region associated with variation in the environment?".

using combining information across tightly linked sites to identify regions of the genome under selection. Typically GEA studies examine patterns of genetic variation across a landscape at many polymorphic sites.

Materials and Methods

The Weighted-Z Analysis

In this study, we propose the Weighted-Z Analysis (hereafter, the WZA) for combining information across linked sites in the context of GEA studies. The weighted-Z test combines *p*-values from multiple independent tests into a single score with each test given a weight that is proportional to the inverse of its error variance (Whitlock 2004). Inspired by Weir and Cockerham's (1984) method for combining estimates of F_{ST} across sites, we use a marker's allele frequency to determine weights when performing the Weighted-Z test on GEA data. At a given polymorphic site, we denote the average frequency of the minor allele across populations as $\bar{p}$ ($\bar{q}$ corresponds to the major allele). The product $\overline{pq}$ provides an estimate of the variance in allele frequencies among populations, so is appropriate as a weight.

We combine information from multiple SNPs present in a particular region into a single weighted-Z score (Z_W). For genomic region k, which contains n polymorphic sites, we calculate

$$Z_{w,k} = \frac{\sum\limits_{i=1}^{n} \overline{p_i q_i} z_i}{\sqrt{\sum\limits_{i=1}^{n} (\overline{p_i q_i})^2}},$$

(1)

where $\overline{p_i}$ and $\overline{q_i}$ are the average allele frequencies across demes for polymorphism i and z_i is the standard normal deviate calculated from the one-sided *p*-value for SNP i. In practice, GEA analyses might return Bayes factors or false-discovery corrected *q*-values. In such cases, z scores calculated from empirical *p*-values may be used.

The Top-Candidate test

Yeaman et al (2016) proposed a method for combining information across sites in genotype-environment association studies that they referred to as the top-candidate test. The top-candidate test attempts to identify regions of the genome involved in local adaptation under the assumption that regions involved in local adaptation may contain multiple sites that exhibit strong correlation with environmental variables. To apply the top-candidate test, outlier SNPs are identified from the genome-wide distribution of test statistic, for example Yeaman et al (2016) classified *p*-values above a the 99th percentile particular percentile threshold genome-wide (we used the 99th percentile) are classified as outliers. The frequency of outliers SNPs in analysis windows is then compared to the total number of SNPs in the window. A binomial test is used to determine whether a given window has an excess of outliers relative to the genome-wide expectation. Analysis windows with a *p*-value less than 0.0001 were taken as "top-candidates" for local adaptation. Note that in Yeaman et al (2016) genes, as well as up and downstream flanking sequence, were used as analysis windows.

Simulating local adaptation

Genotype-environment association studies are often performed on large spatially extended populations. However, it is computationally infeasible to model selection and linkage in large chromosome in such a populations, so we scaled population genetic parameters to tractably model large populations. Table **??** shows the parameters of the organism that we evolved *in silico*. In the Appendix, we outline and justify the approach we used to scale relevant population genetic parameters in our simulations of local adaptation. All simulations were performed in SLiM v3.4 (Messer and Haller 2018).

We simulated genomes with a single chromosome containing 1,000 "genes". Each "gene" was 9,999bp long and recombined at a rate of $r = 10^{-7}$, between each gene a single base-pair recombined at a rate of $r = 0.005$. Thus, our simulated chromosomes were 10Mbp long each, but modelled a 599 cM chromosome.

The simulated populations inhabited a 14×14 2-dimensional stepping-stone population.

We simulated three kinds of environment. The first was a reduced representation of climatic variation across British Columbia, Canada. We downloaded the map of degree days greater than 0 (DD0) for British Columbia from ClimateBC (website; REF). From the DD0 map (Figure 1A), we extracted the data for a 99x99 grid using Dog Mountain, BC as the reference point in the South-West corner. We divided this map into a 14x14 grid. Each cell corresponded to an area of XXkm^2. We calculated the mean DD0 for each cell in the grid. We then converted the mean DD0 scores into Z-scores and rounded values up to the nearest third. These data were then used as the phenotypic optima for population models in SLiM. The data from the BC Map were then ordered along one axis of the 2D stepping-stone, we refer to this map as the *cline* map (Figure 1C). Finally, we truncated the distribution of Z-scores from the BC map at +3, setting all demes with a DD0 Z score greater than or equal to 3 to 3, and all others to -1. We refer to this map as the *Truncated map* (Figure 1D).

In addition, we simulated local adaptation in a metapopultion with no spatial structure (i.e. an island model). The first is the island model, which represents an unstructured metapopulation.

We used the standard expression for Gaussian stabilising selection,

$$W(z_{i,j}) = exp\left[\frac{-(z_{i,j} - \theta_j)^2}{2V_s}\right],$$

where z_i is the phenotype of the i^{th} individual in environment j, θ_j is the phenotypic optimum of environment j, and V_s is the variance of the Gaussian fitness function.

To test the performance of the weighted-Z analysis (WZA), we modelled populations adapting to various environments. Figure **??** shows a diagram of each of the populations simulated.

Covariance of phenotype and environment

In our simulations, we used the covariance between and environment as a measure of a gene's relevance for local adaptation. We calculated the average phenotypic effect of each gene as follows. We then used $Cov(PB_g,$

env)/Cov(*PB*, *env*) as a measure of a gene's contribution to local adaptation.

In our simulations, phenotypic variance (σ_P^2) was generated solely by genotypes, i.e. there were no environmental effects. Local adaptation generates variance in phenotypes between populations (σ_{PB}^2). As described above, the simulations incorporated a stochastic mutation model, so from replicate to replicate the effect size of alleles and their locations in the genome varied. As a result, the genes that contributed to local adaptation varied across simulation replicates. We therefore determined the contribution each gene made to local adaptation by calculating the proportion of phenotypic variance among populations explained by the SNPs in each gene. For each gene that contributes to phenotypic variation there are k causal SNPs each with a phenotypic effect of α_k. We use $\boldsymbol{v_g}$ to refer to the column vector of phenotypic effects for each of the k causal SNPs in gene g. In each population there are n diploid individuals and we have $\boldsymbol{M_d}$, an $n \times k$ matrix in which the genotype of each individual at each causal SNP is coded as 0, 1 or 2 corresponding to aa, aA and AA genotypes, respectively. The contribution that each gene makes to the overall phenotype in each population is calculated as $C_{g,d} = \sum \boldsymbol{M_{g,d}} \boldsymbol{v_{g,d}}$. The variance in $C_{g,d}$ gives us a measure of the phenotypic variance between popuatlions generated by each gene ($\sigma_{PB,g}^2$). We then calculate the proportion of variance explained by each gene (PVE_g) as $\sigma_{PB,g}^2 / \sigma_{PB}^2$

$$PVE_g = \frac{\sigma_{PB,g}^2}{\sigma_{PB}^2}. \tag{2}$$

Note that PVE_g does not provide a measure of local adaptation, merely a measure of how much phenotypic variation between populations can be explained by a a particular gene.

We recorded the Cov(fitness, environment) as a measure of a gene's relevance to local adaptation in a particular simulation replicate. The three upstream and three downstream genes.

Analysis of simulation data

We added neutral mutations to each simulated tree sequence at a rate of 1×10^{-8}. While this gave us a population scaled mutation rate of $4N_e d\mu = 0.00078$, and resulted in an average of 30 SNPs per gene that passed a minor allele frequency filter of 0.05.

For each SNP, we calculated the allele freuqncy for each deme. We calculated Kendall's τ between allele frequncies and the environment for that population. We chose Kendall's τ as an uncorrected GEA statistic as it can handle ties in data.

Tree sequences were manipulated using the tskit package. Mutations were added to trees using msprime (REF) through the PySLiM package (version). F_{ST} and r^2 (a measure of linkage disequilibrium) were calculated using custom Python scripts that invoked the scikit-allel package (REF).

Detecting outlier regions using a SNP-based approach

We performed an individual SNP-based approach to identify genes involved in local adaptation. The scores for individual SNPs across a simulation replicate were ranked. The gene containing the SNP with the most extreme test statistic (e.g. the smallest p-value) was scored as a hit and all other SNPs in the identified gene were subsequently ignored. The second-most extreme test statistic SNPs present in that gene were then ignored.

wih the individually highest test-statistic was then When analysing genome-wide SNP data, a set of highly significant points may be in tight linkage (i.e. there may be multiple outliers).

Analysis of data from Lodgepole pine

We re-analysed a population genomic dataset collected for lodgepole pine distributed across the North West of North America. The data were initially generated and described by Yeaman et al (2016). Initially, the top-candidate test was applied to this data. We calculated Z_W scores for the same genes analysed by Yeaman et al (2016). Data were accessed from the Dryad respository associated with Yeaman et al (2016) (DRYADLINK)

Data Availability

The simulation configuration files and code to perform the analysis of simulated data and generate the associated plots are available at github/TBooker/GEA. Tree-sequence files for the simulated populations are available at Dryad and all processed GEA files are available on ().

Results

Simulations of local adaptation

Our stepping-stone populations exxhibed strong signals of isolation-by-distance, F_{ST} was greatest between physically distant genes.

Patterns of linkage disequilibrium

A
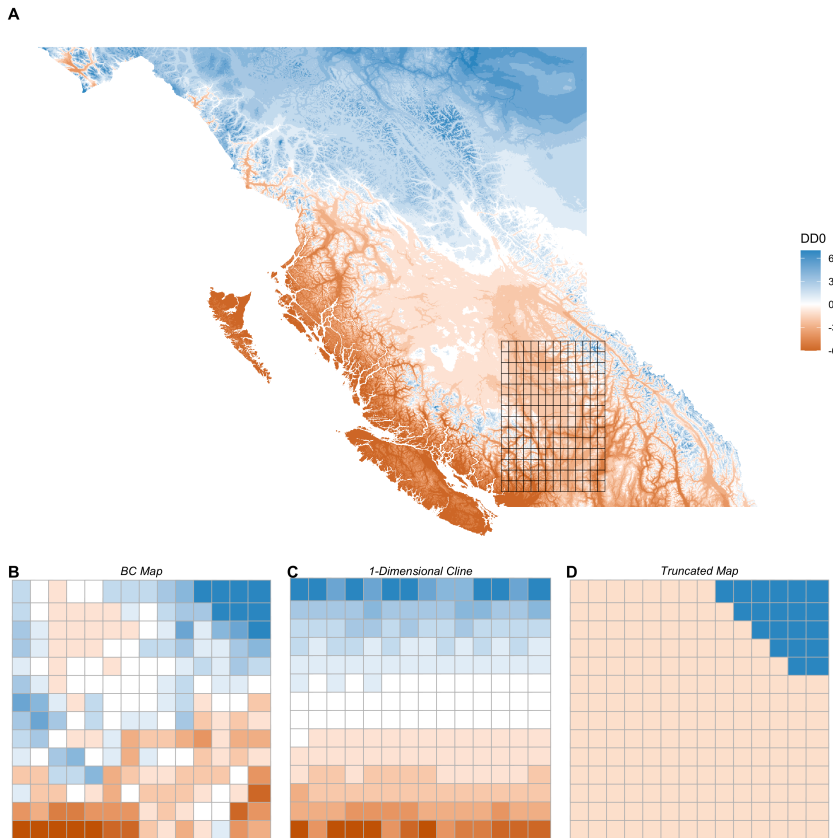
Figure 1 A) Degree days below zero across British Columbia, the inset grid in A shows the locations we used to construct phenotypes for our simulated populations. B) A highly discretized map of DD0 in Southern British Columbia, we refer to the map in B as the BC map. C) A 1-dimensional gradient in phenotypic optimum, we refer to this as the 1D gradient map. D) A heterogenous distribution of phenotyic optima.

A Island Model – Uncorrected

B BC Map – Uncorrected
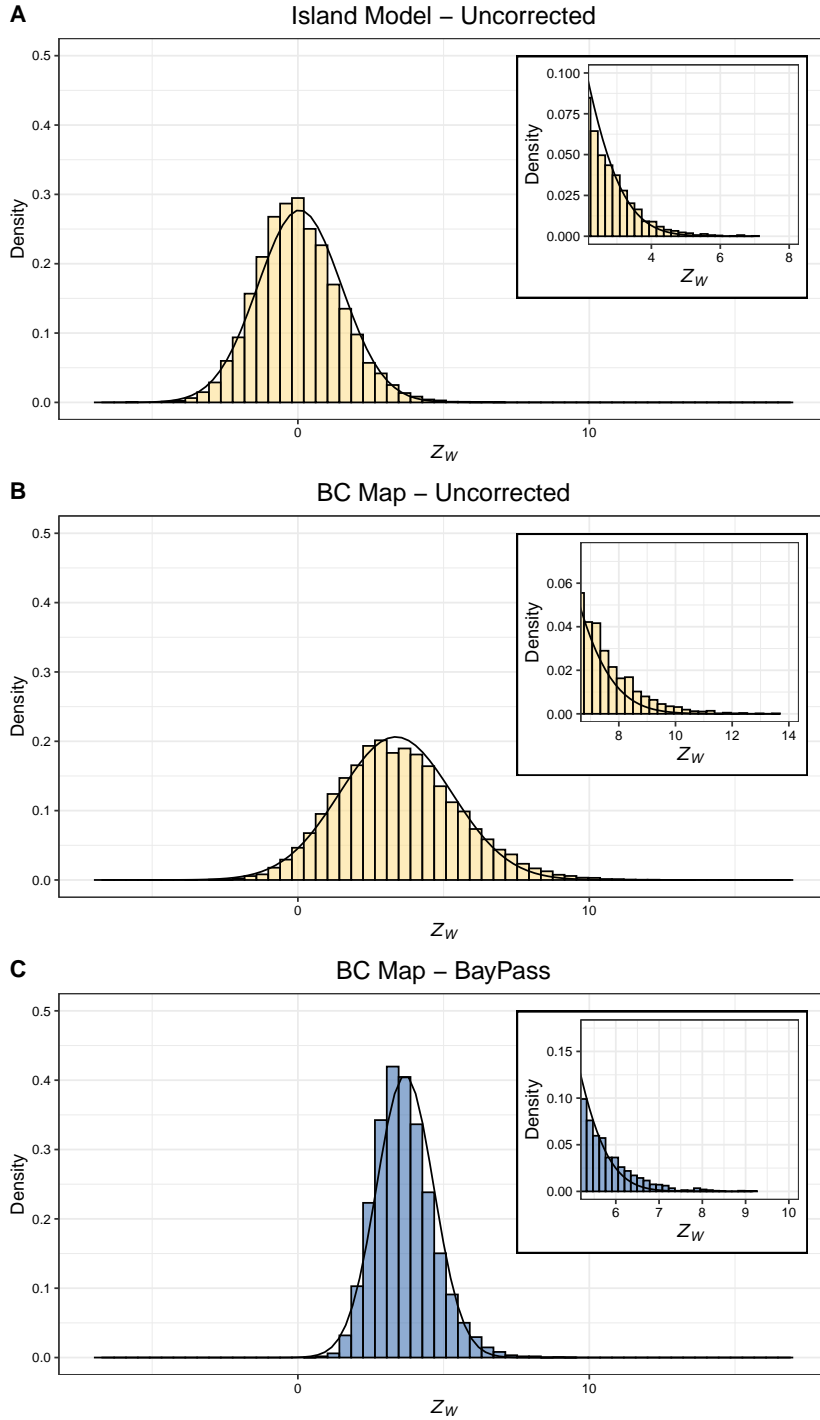

C BC Map – BayPass

Figure 2 Density histograms of WZA scores for three cases. A) An finite island population, Z_W scores were calculated using uncorrected p-values obtained from Kendall's τ. B) The BC-map, Z_W scores were calculated using uncorrected p-values obtained from Kendall's τ. C). Inset in each panel is a histogram focussing on the upper tail of the Z_W distribution.

To assess the statistical properties of the WZA and the top-candidate test, we first performed GEA analyses on populations structured according to an island model. While highly unrealistic, analysing this model allowed us to determine the statistical properties of the WZA and the top-candidate test without the need to correct for

the confounding effects of population structure.

The distribution of Z_W scores obtained for populations evolving under strict neutrality was very close to the expectation of the standard normal distribution. The mean Z_W was 0.00X and the variance was 1.XXX. Figure **??**A shows the distribution of Z_W scores obtained when analysing a sample consisting of 50 individuals from 40 demes (2,000 total). However, simulations modelling local adaptation in the island model resulted in a skewed distribution of Z_W scores for neutral genes. Figure **??**C shows that adaptation elsewhere in the genome can generate a background level of correlation with the environment, causing the mean Z_W to be greater than 0. Indeed, the mean Z_W was 0.00X and the variance was 1.XXX in this case. This isolation-by-adaptation means that it is not possible to convert Z_W scores to parametric p-values. Some degree of isolation-by-adaptation should probably be expected in natural organisms

The populations we simulated had **5** chromosomes, one of which was strictly neutral. Applying the WZA to simulated genes from a neutrally evolving chromosome in a locally adapted population allows us to test whether isolation-by-adaptation causes a

*Comparison of the WZA and the top-candidate test*

HELLO

Figure 3 The proportion of true positives detected across a specified search effort when a model of local adaptation by directional selection was assumed.

Figure 4 The proportion of true positives detected across a specified search effort when a model of local adaptation by stabilising selection was assumed.

Figure 5 The proportion of true positives recovered when the measured environment is correlated with selection pressure to varying degrees. The correlation between environment and selection pressure is shown above the plots. Results are shown assuming the BC Map and the model of stabilising selection.

Recombination rate variation

Recombination rates vary widely among taxa but also within the genome (Stapley et al REVIEW). Genomic regions of low recombination exhibit greater variance in population genetic summary statistics than do more highly recombining ones, complicating statistical inference. Analysing genomic data using windows of a fixed physical size in genomes with varying recombination rates can lead to false positives in regions of low recombination rate (Booker et al 2020). Alternatively, studies might choose to focus on genes rather than analysis windows. In addition, analysis of real data might be focussed on genes, rather than on In our simulations, all gene experienced the same recombination rate, though that is highly unrealistic, we did so for statistical convenience.

Figure 6 The distribution of Z_W scores under different recombination rates

Application of the WZA to data from lodgepole pine

To demonstrate how one might apply the WZA to the analysis of real data, we re-analysed the lodgepole pine (*Pinus contorta*) data from Yeaman et al (2016). Briefly, Yeaman et al (2016) collected samples from 666 populations across British Columbia and Alberta, Canada and from Northern Washington, USA. Individuals in each Data were downloaded from the Dryad repository associated with the paper.
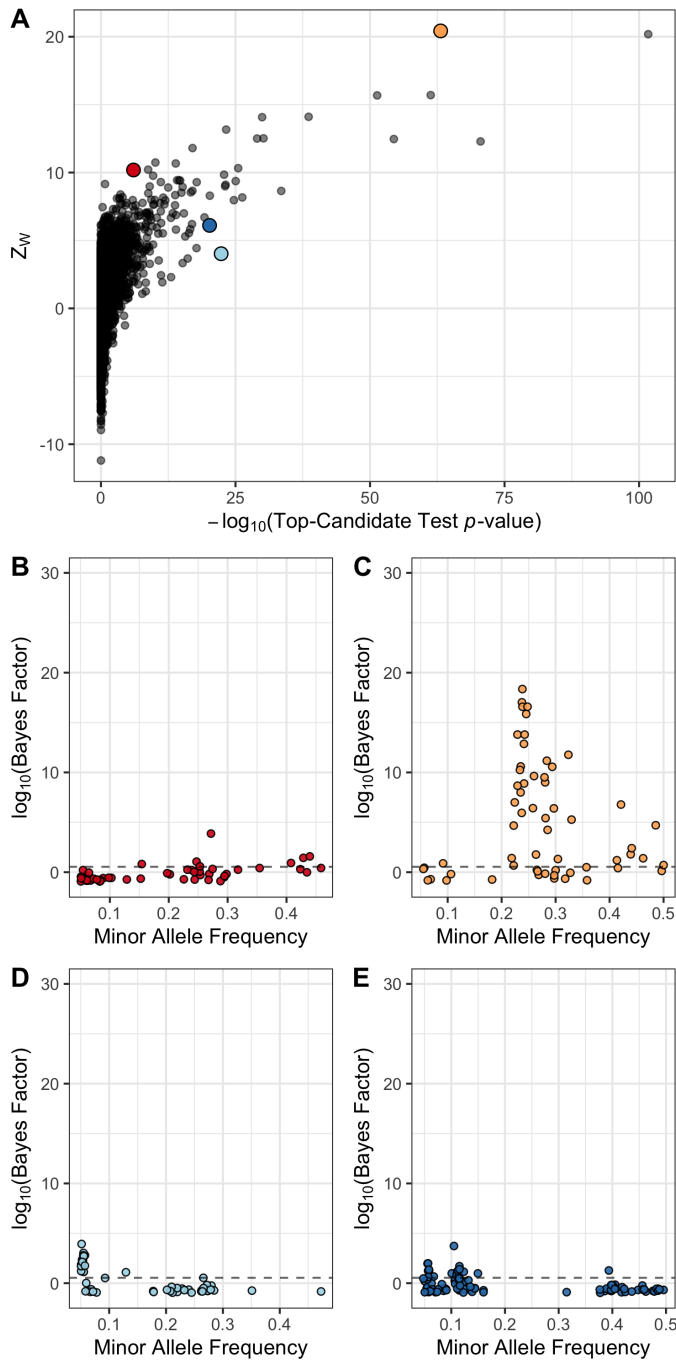
Figure 7 The WZA applied to GEA results Lodgepole Pine for degree days below 0 (DD0). A) Z_w scores compared to scores from the top-candidate test for each of the genes analyesd by Yeaman et al (2016). Panels B-E show the results for $-log_{10}(p$-values) for Spearman's ρ applied to individual SNPs against minor allele frequency (MAF). The dashed horizontal line in B-D indicates the 99^{th} percentile of GEA $-log_{10}(p$-values) genome-wide.

Discussion

Paragraph about weak selection... Populations may be adapted to

There are philosophical reasons as to why the WZA should be preferred. First, the top-candidate test assumes that there is a fraction of the genetic markers analysed that are tagging causal variants (i.e. that there are true positives in the dataset). This is undesirable, because there may well be no detectable variation that

contributes to local adaptation present, i.e. genuine genotype-environment correlations may be very weak and the study are simply underpowered. Secondly, the top-candidate test gives equal weight to all markers. However, alleles at different frequencies possess different levels of information about population history. A final related point is that all SNPs that have exceeded the significance threshold are treated identically. For example, with a significance threshold of 0.01, genomic regions with only a single outlier are treated in the same way whether that outlier has a p-value of 0.009 or 10^{-10}.

Population expansions can cause allelic surfing, where regions of the genome "surf" to high frequency on the leading edge of the expanding population. This can leave heterogeneous patterns of linkage disqequilibirum in the genome (Excoffier et al). Indeed, such allelic surfing can resemble selective sweeps of strongly beneficial mutations (Ref). When population edges experience environmtal conditions that are highly dissimilar to the rest of a species' range, allelic surfing could generate a spurious signal of genotype environment correlation.

When analysing real datasets, researchers should be mindful that analysis windows of a constant physical size may generate statistical artefacts as we outlined in our previous study (Booker et al 2020).

Acknowledgements

Bibliography

Appendix

Consider a hypothetical metapopulation of 1 million individuals distributed evenly among 196 demes. It would be computationally intractable to simulate all 10^6 individuals forward-in-time, incorporating adaptation to environmental heterogeneity across a landscape and recombining chromosomes. We scaled several population genetic parameters to model a large population by simulating a much smaller population. In the following sections, we outline and justify the approach we used to scale pertinent population genetic parameters.

Recombination rates In panmictic populations, linkage disequilibrium can be predicted by the scaled recombination parameter $\rho = 4N_e r$, where r is the recombination rate per base-pair and N_e is the effective population size. In structured populations, LD is elevated above the panmictic expectation and can be described by the effective size of the local population (or deme), the recombination rate and the migration rate (McVean paper). Assuming a recombination rate of 1 cM/Mbp, the hypothetical organism would have $4N_d r = 0.0002$. To achieve levels of LD-decay in our simulations that are similar to those expected in our hypothetical organism, we set $4N_d r = 0.0002$, but with only 100 individuals per deme that gave a per base pair recombination of 5.10×10^{-7}.

Selection coefficients It is difficult to choose a realistic set of selection parameters for modelling local adaptation because there are, at present, few estimates of the distribution of fitness effects for mutations that have spatially divergent effects. However, common garden studies of a variety of taxa have estimated fitness differences of up to 50% between populations grown in home-like conditions versus away-like conditions (REF). Motivated by such studies, we chose to parameterise selection using the maximum possible fitness difference between home versus away environments. By setting the maximum reduction to ! have demonstrated in a variety of taxa that there

Migration rate For the migration rate, we worked backwards. We set out to achieve F_{ST} across the metapopulation of around 0.05, as has been reported for a numerous widely distributed tree species (REF). For an island model, we can

Mutation rate

Table S1 Population genetic parameters of a hypothetical organism, and how they are scaled in the simulations. The meta-population inhabits a 14×14 2-dimensional stepping stone. Parameters are shown for a population with 12 loci subject to directional selection.

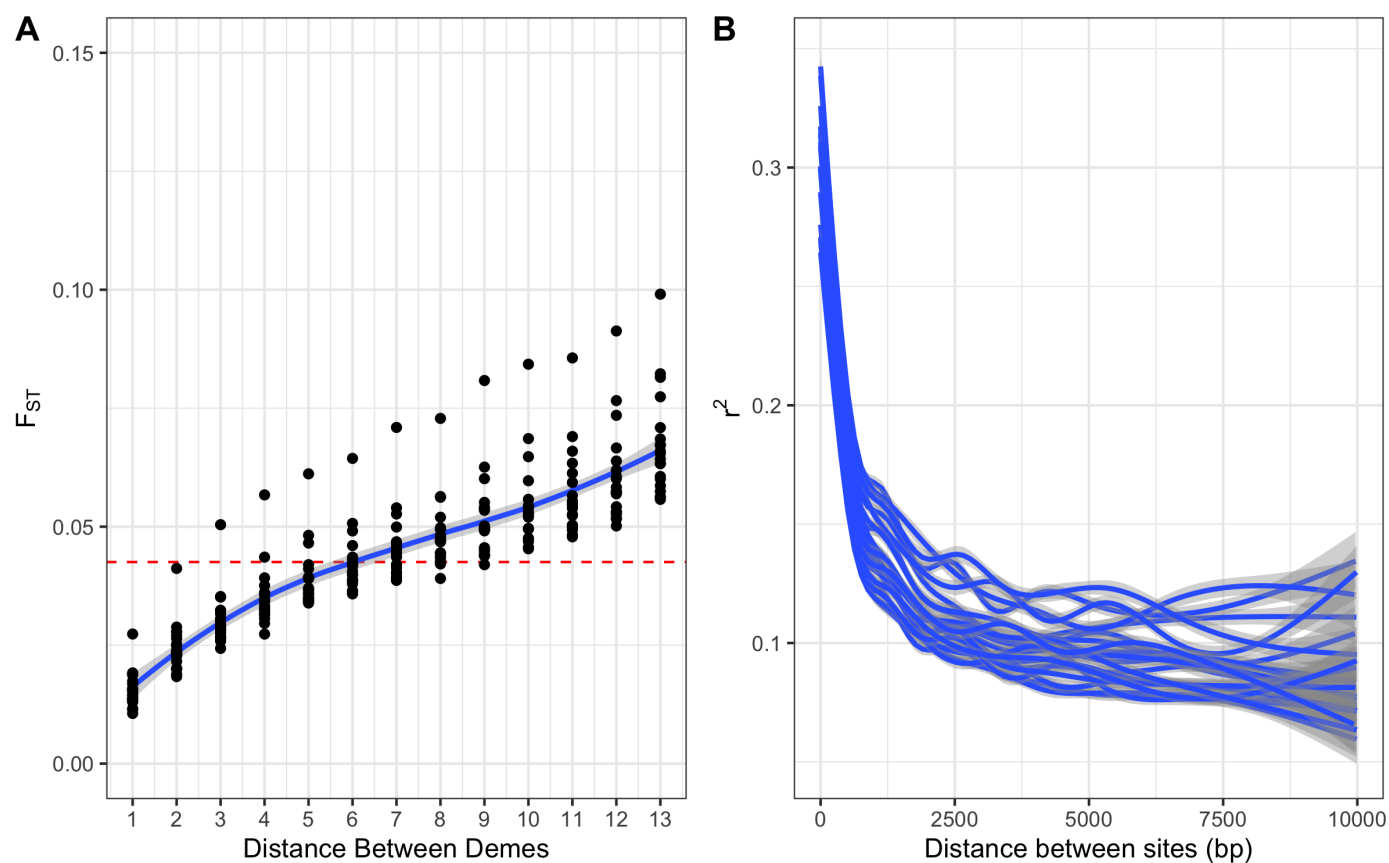| Parameter	Hypothetical Biological Value	Scaled Parameter	Unscaled (simulation)
Global population size (N_e)	10^6	-	19,600
Number of demes (d)	196	-	196
Local population size (N_d)	5,100	-	100
Recombination rate (r)	10^{-8}	$4N_d r = 2.04 \times 10^{-4}$	5.10×10^{-7}
Selection coefficient (s_{Max})	0.0001	$2N_d s_{Max} = 0.6$	0.003
Migration rate (m)	9.80×10^{-4}	$2N_d m = 10$	0.05
Functional mutation rate (μ_α)	2×10^{-10}	$4N_e \mu_\alpha = 0.0008$	10^{-8}

Figure S1 Summary statistics from simulations. A) shows the F_{ST} between pairs of demes in stepping-stone populations, the average across replicates is . B) shows LOESS smoothed LD, as measured by r^2, between pairs of SNPs, each line corresponds to a single simulation replicate.
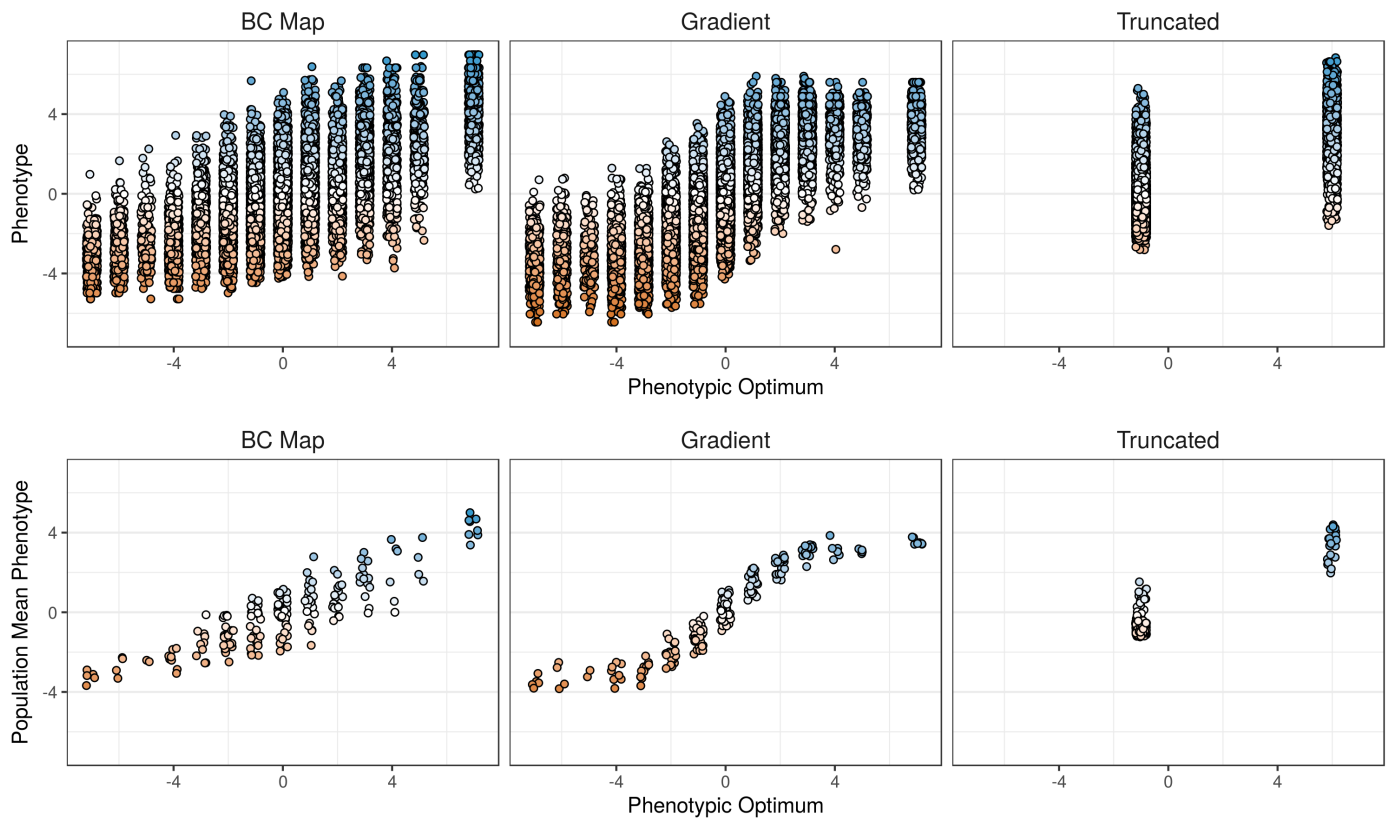
Figure S2 Individual and population mean phenotypes observed in representative simulations for each of the environment maps simulated. A small amount of horizontal jitter was added to points for visualisation purposes
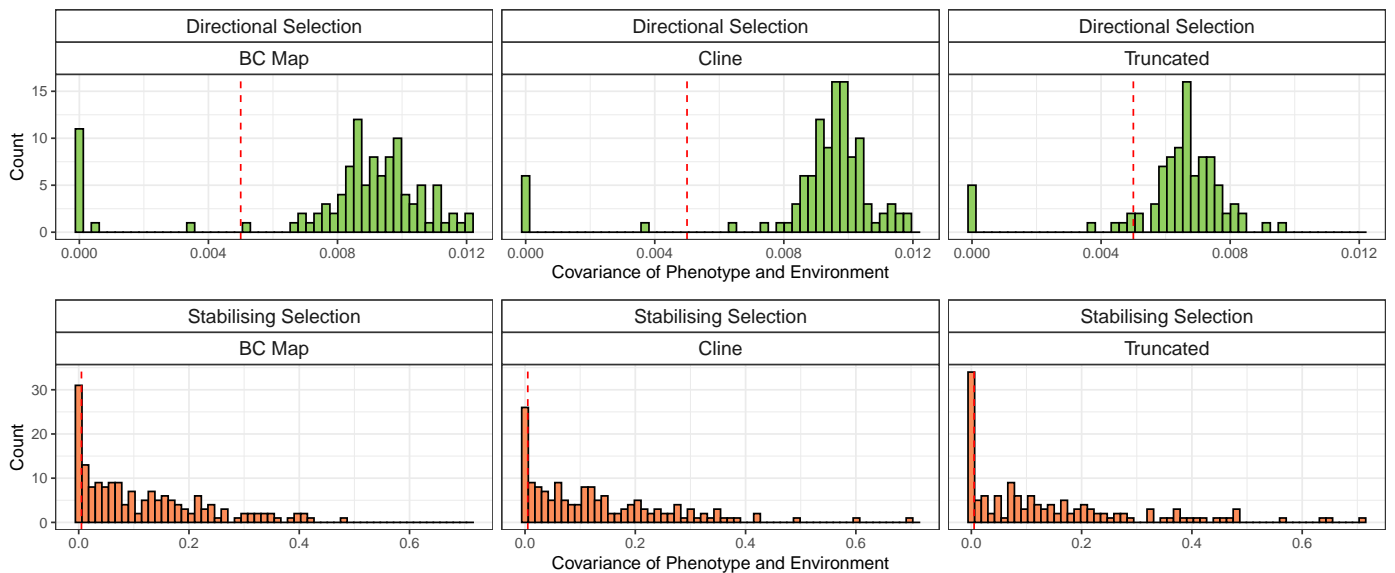

Figure S3 The effect size distribution from simulations of local adaptation. **FIX X-AXIS LABELS**
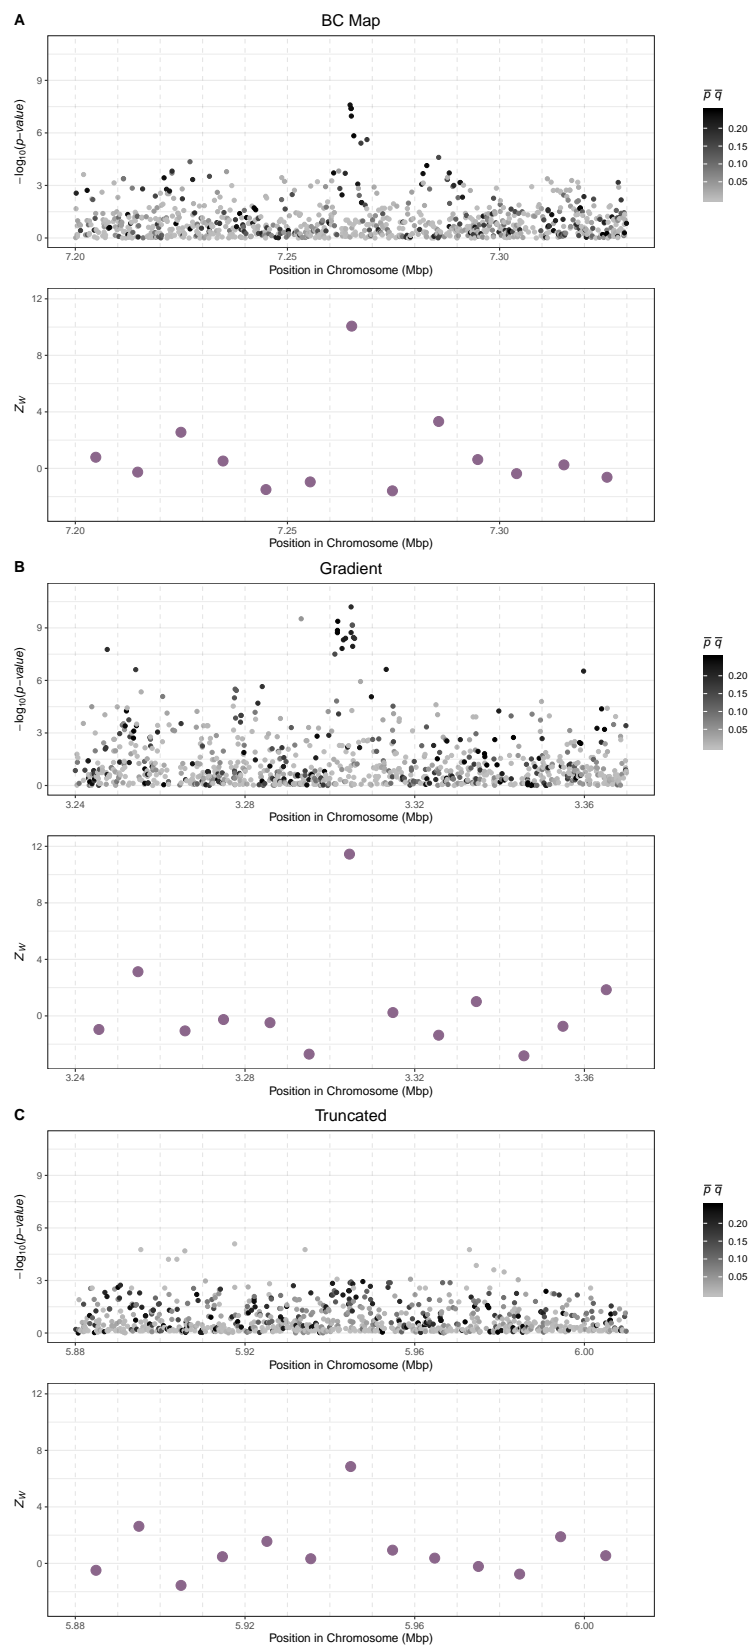
Figure S4 The effect size distribution from simulations of local adaptation. **FIX X-AXIS LABELS**

Figure S5 The effect size distribution from simulations of local adaptation. **FIX X-AXIS LABELS**

Chromosomal architecture